

2.

Task 2



Fig.1 imshow

3.

Task 3




Fig.2 convert to grayscale
with `rbg2gray()`

4.

Task 4




Fig.3 convert grayscale
image to double using `im2double()`

Task 4 double




Fig.4 rescaled result
after using `double()` to convert

- What is the difference between these two conversions?

`im2double()`:

converts the intensity of our grayscale image to double precision and rescales the result to the range `[0 1]`, which is the default display range for `imshow()`.

`double()`:

converts the grayscale image to double precision in range `[0 255]`, which will display a blank image when using `imshow`.

So in order to preserve the image intensity, we need to rescale the image back to range `[0 1]` by dividing it with 255, which yields Fig.4.

- What is the difference between a double image and an unsigned integer image?

unsigned integers values are in a fixed range, eg: [0 128], [0 255], and each pixel can only adopt one integer value as its intensity level.

double images typically range from 0 to 1, but the number of values each pixel can adopt would be the same as the double precision.

5.

Implementation: Extracting submatrix from image by matrix indexing

Task 5




Fig.5 obtained by `Crop(img,5,5,50,70)`

6.

Implementation: use nested for-loops to implement 2D convolution and convolve the image intensity with a nxn averaging window (all values are 1s).

- Window size = 3:




Fig.6 my implementation
using smaller windows on pixels that
have fewer than $n \times n - 1$ neighbors

- Window size = 15



Fig.8 my implementation




Fig.7 using conv2()
zero padding results in dark
borders




Fig.9 using conv2()
the effect of zero-padding
increases with window size